

# Technical Neural Networks Topics:

## 1. Basis

- Biological Background
- Learning Paradigms
- History of Neural Networks
- Perceptron

## Neural Network Paradigms

→ For every paradigm Structure, Functionality and Learning Rule has to be learned

## 2. Supervised Learning

2.1 MLP

2.2 RBF

2.3 ART

2.4 Neocognitron

2.5 SVM

2.6 Recurrent MLPs

- Elman-Network

- Jordan-Network

## 4. Reinforcement Learning

- Reward Functions

- Value Functions

- Learning a policy

## 3. Unsupervised Learning

3.1 SOM

3.2 Multi-SOM

3.3 Neural-Gas

3.4 Multi-Neural-Gas

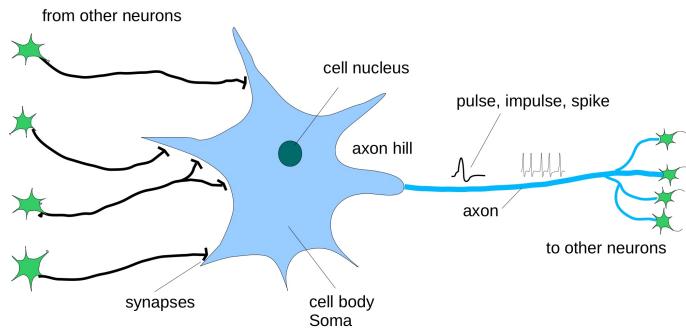
3.5 k-means clustering

3.6 Learning Vector Quantization

# I. Basis

## Biological Background

### Biological Neurons:



### Principles:

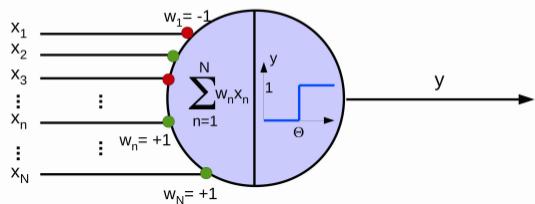
- High interconnectivity
- Neurons structured in layers
- Neurons are activated above a certain threshold

### Hebbian Learning Rule:

Connection between neurons is strengthened if both neurons are activated in the same time window

$$\Delta W_{nm} = \eta * y_m * x_n$$

### Technical Neuron:



## History of Technical Neural Networks

### 1940-1960: Beginning

→ first neural paradigms such as the Perceptron inspired by biological background, the applicability is not yet proven

### 1960-1982: "Quiet Years"

→ no funding and progress in the area due to no applicable learning rule, a single perceptron can only deal with linear separable data

### 1980-1990: Renaissance

→ the quiet years result in many different neural paradigms leading to a restart and explosion in ideas

### 1990-1999: Consolidation

→ Neural networks are used in many different areas through efficient learning algorithms

### Since 2000: 2<sup>nd</sup> Start-up

→ Through new paradigms and higher computational power neural networks can solve highly complex tasks

## Learning Paradigms:

### Supervised Learning:

Learning is performed through the utilization of a teacher value  $\hat{Y}$ . Given input  $X$ , the error between the output  $Y$  and  $\hat{Y}$  is trained to be minimized.

### Unsupervised Learning:

Only data driven learning. The learner tries to map statistical properties of the input  $X$ .

### Reinforcement Learning:

The learner only receives a sparse feedback to his output  $Y$  which he evaluates his action on.

## Applications of Neural Networks

### Function Approximation:

A neural network can be used as a universal function approximator. The goal is to approximate a goal function as closely as possible given samples from it.

### Pattern Recognition:

The goal is to map the input to a discrete set of values which classifies the input to pre-defined or self-learned classes.

## Binary Classification

Task: Assign to an input a binary class value

Further assumption: linear separable data

Goal of Learning:

Find parameters  $\alpha$  of function/classifier  $f_\alpha$ , such that:

1. Minimize Risk:  $R(\alpha) = \int \frac{1}{2} |y - f_\alpha(\vec{x})| dP(\vec{X}, y)$

→ usually not observable

Empirical Risk:  $E_\alpha = \frac{1}{L} \sum_{i=1}^L \frac{1}{2} |y - f_\alpha(\vec{x}_i)|$

→ Training error

→ Approximation of the risk, can deviate due to overfitting etc.

## 2. Structural Risk Minimization

Minimize the gap between the risk and the empirical risk

$$R(\alpha) \leq E_\alpha + \Theta(h, L, \eta)$$

Confidence Term  $\Theta$ :

$$\Theta(h, L, \eta) = \sqrt{\frac{h \left( \log\left(\frac{2L}{h}\right) + 1 \right) - \log\left(\frac{\eta}{4}\right)}{L}}$$

$L$ : Number of training patterns

$h$ : VC capacity of function class  $f_\alpha$

$\eta$ : Probability

For  $h < L$ : This inequality holds with  $(1-\eta)$

Reduce  $\Theta$ :

1. Take a larger amount of training data  $L$
2. Keep VC-dimension low, prevent overfitting

VC Capacity  $h$ : Measure to quantify the expressive power of a class of functions

Maximal number of points that can be shattered by the class of functions.

Shattering: All dichotomies of  $h$  points can be realized

1. Show there exist a set of  $h$  points that can be shattered
2. Show there exist no set of  $(h+1)$  points

Hyperplane

$$\left\{ \vec{Z} \mid (\vec{W} \cdot \vec{Z}) + b = 0 \right\}$$

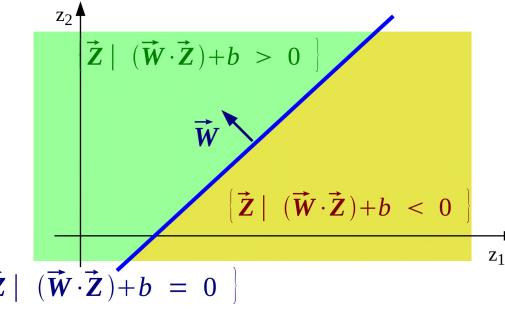
In 2D:

$$(\vec{W} \cdot \vec{Z}) + b = 0$$

$$w_1 z_1 + w_2 z_2 + b = 0$$

$$w_2 z_2 = -w_1 z_1 - b$$

$$z_2 = -\frac{w_1}{w_2} z_1 - \frac{b}{w_2} \quad \text{case } w_2 \neq 0$$



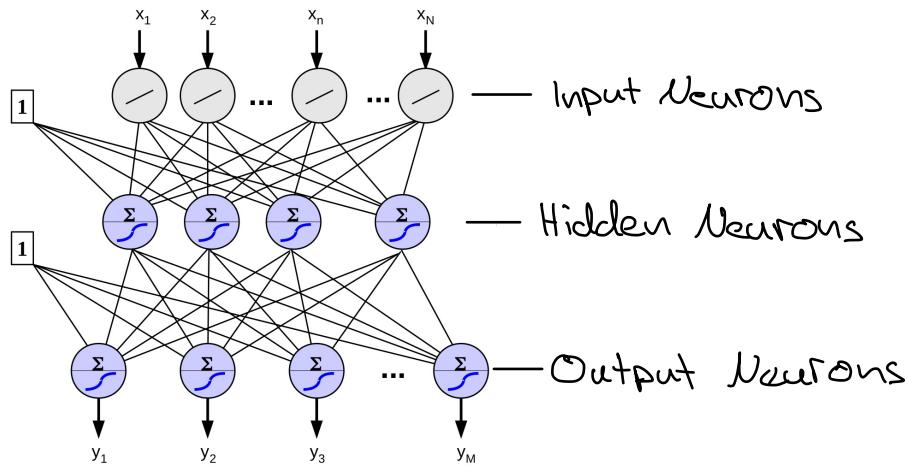
Binary classifier:

$$f_{W,b}(\vec{Z}) = \text{sign}( (\vec{W} \cdot \vec{Z}) + b )$$

## 2. Supervised Learning

### Multi-Layer Perceptron

#### Structure:



#### Learning Rule: Backpropagation of Error

##### Error Function:

$${}^p E = \frac{1}{2} \sum_{m=1}^M ({}^p \hat{y}_m - {}^p y_m)^2$$

##### Weight Change:

$$\Delta w_{ij} = \eta \cdot \delta_j \cdot \tilde{out}_i$$

\Delta\text{-rule}

##### Output Neuron:

$$\delta_m = (\hat{y}_m - y_m) \cdot f'(\text{net}_m)$$

##### Hidden Neuron:

$$\delta_h = \left( \sum_{k=1}^K \delta_k w_{hk} \right) \cdot f'(\text{net}_h)$$

#### Derivation:

Calculate the gradient w.r.t the error function to minimize the error through gradient descent

##### Output Neuron:

$$\Delta w_{hm} = -\eta \frac{\partial {}^p E(w_{hm})}{\partial w_{hm}}$$

| Definition

$$\frac{\partial {}^p E}{\partial \text{net}_m} \cdot \frac{\partial \text{net}_m}{\partial w_{hm}}$$

| Chain rule

$$\frac{\partial {}^p E}{\partial {}^p y_m} \cdot \frac{\partial f(\text{net}_m)}{\partial \text{net}_m}$$

| Chain rule

$$\delta_m = ({}^p \hat{y}_m - {}^p y_m) \cdot f'({}^p \text{net}_m)$$

$$\Delta w_{hm} = \eta \cdot ({}^p \hat{y}_m - {}^p y_m) \cdot f'({}^p \text{net}_m) \cdot \tilde{out}_m$$

## Hidden Neurons:

$$\Delta w_{hm} = -\eta \frac{\partial E(w_{hm})}{\partial w_{hm}}$$

$$\frac{\partial E(w_{gh})}{\partial p_{net_h}} \cdot \frac{\partial p_{net_h}}{\partial w_{gh}} \quad | \text{ chain rule}$$

$$-\delta_n = \frac{\partial E}{\partial \text{out}_n} \cdot \frac{\partial \text{out}_n}{\partial p_{net_h}} \cdot \frac{\partial f(p_{net_h})}{\partial p_{net}}$$

$$\sum_{k=1}^K \frac{\partial E}{\partial p_{net_k}} \cdot \frac{\partial p_{net_k}}{\partial \text{out}_n}$$

$$\frac{\partial (\sum_{j=0}^H \text{out}_j w_{j,k})}{\partial w_{jk}}$$

$$- \sum_{k=1}^K \delta_k \cdot w_{hk}$$

$$\delta_n = \left( \sum_{k=1}^K \delta_k \cdot w_{hk} \right) \cdot f'(p_{net_h})$$

$$\Delta w_{hm} = \eta \cdot \left( \sum_{k=1}^K \delta_k \cdot w_{hk} \right) \cdot f'(p_{net_h}) \cdot p_{out_g}$$

| only input g contributes

| chain rule

| only output h contributes

| contributes of h in next layer

## Transfer Functions:

### Function:

Logistic

$$f_{logist}(z) = \frac{1}{1+e^{-z}}$$

Tanh

$$f_{tanh}(z) = \tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

Relu

$$f_r(z) = \max(0, z)$$

Softplus

$$f_{softplus}(z) = \ln(1+e^z)$$

### Derivation:

$$f'_{logist}(z) = f_{logist}(z) \cdot (1.0 - f_{logist}(z))$$

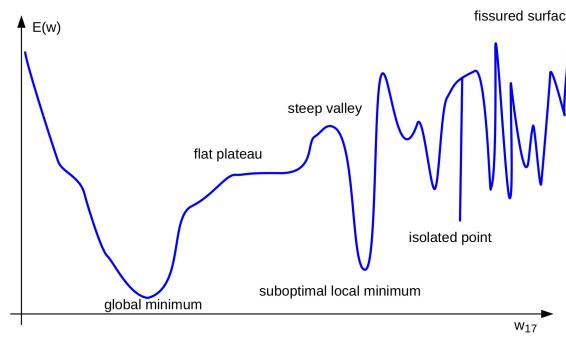
$$f'_{tanh}(z) = \tanh'(z) = 1.0 - \tanh^2(z)$$

$$f'_r = \begin{cases} 1 & , \text{for } z > 0 \\ 0 & , \text{else} \end{cases}$$

$$f'_{softplus}(z) = \frac{1}{1+e^{-z}}$$

## Implementation Details:

Network Choice:  $\frac{\# \text{ Patterns}}{\# \text{ Neurons}} \approx 3$

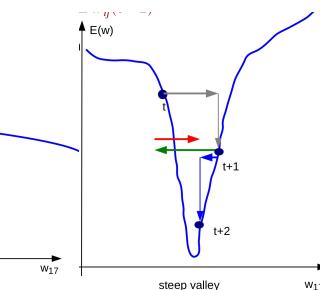
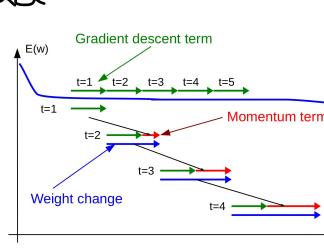


## Training Procedure:

1. Single step learning: Weight update for single patterns
- Batch Learning: Weight update on a basis of a batch of patterns
2. Present patterns in a new random sequence every epoch
3. Momentum: Taking the momentum from the last step helps dealing with plateaus or isolated local minima

$$\Delta w_{ij}(t) = \eta \cdot \delta_j \cdot {}^p\overline{\text{out}}_i + \alpha \cdot \Delta w_{ij}(t-1)$$

Weight change      Gradient descent term      Momentum term



4. Early Stopping: Stop training when the model starts overfitting or progress stagnates

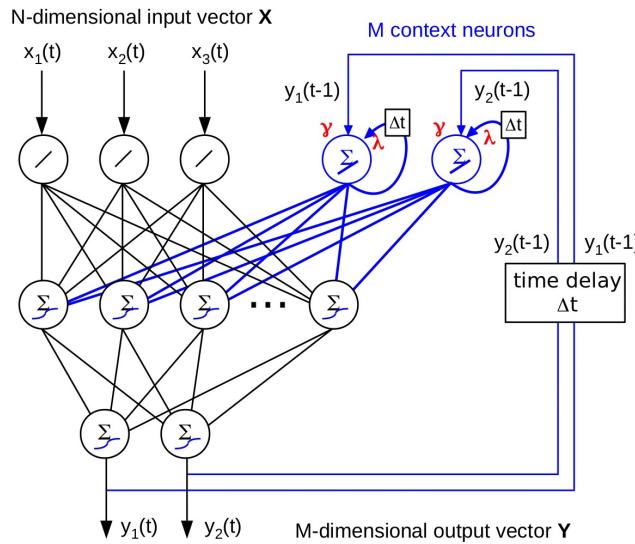
## Recurrent MLPs

Based on normal MLPs with extra feedback connections. This introduces an explicit time delay. The result is a dynamic behavior. Static weights and input generates a sequence of outputs.

## Jordan Networks

Network with additional context neurons that feedback the complete output of the network. The context neurons can have an additional direct feedback to themselves.

### Structure:

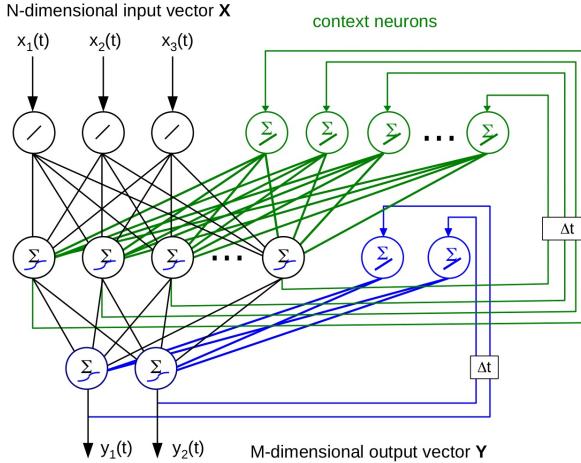


Typically:  $\gamma=1, \lambda=0$

## Elman Networks

Network with feedback connections from every layer to the very same layer

### Structure:



## Operation Modes:

Mode 1: Input is constant  $\rightarrow$  Output converges



Mode 2: Input varying in time  $\rightarrow$  Output converges



Mode 3: Input is constant  $\rightarrow$  Output generates a sequence of states, values or continuous time series



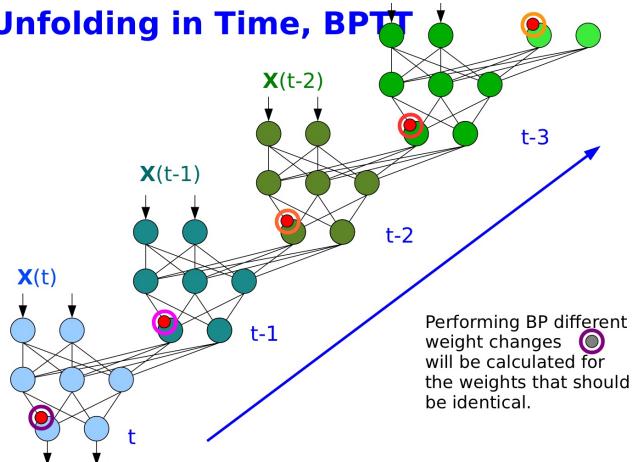
Mode 4: Input varying in time  $\rightarrow$  Output generates a sequence of states, values or continuous time series



## Backpropagation Through Time

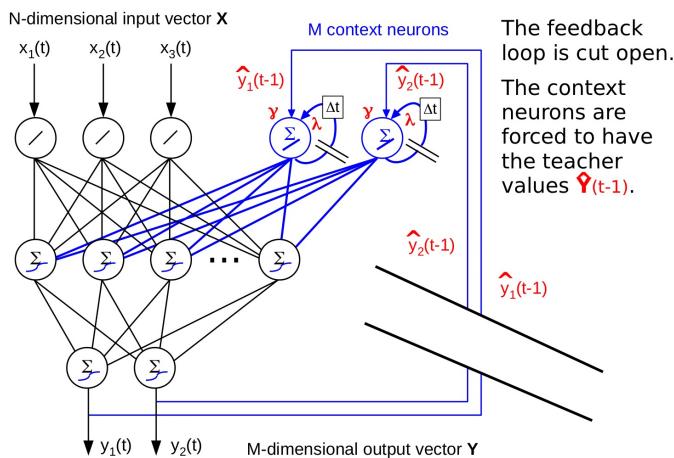
The network is unfolded over time for a fixed number of timesteps. Backpropagation is applied to each unfolded network and the weight changes are accumulated.  
→ high computational/memory efforts

### Unfolding in Time, BPTT



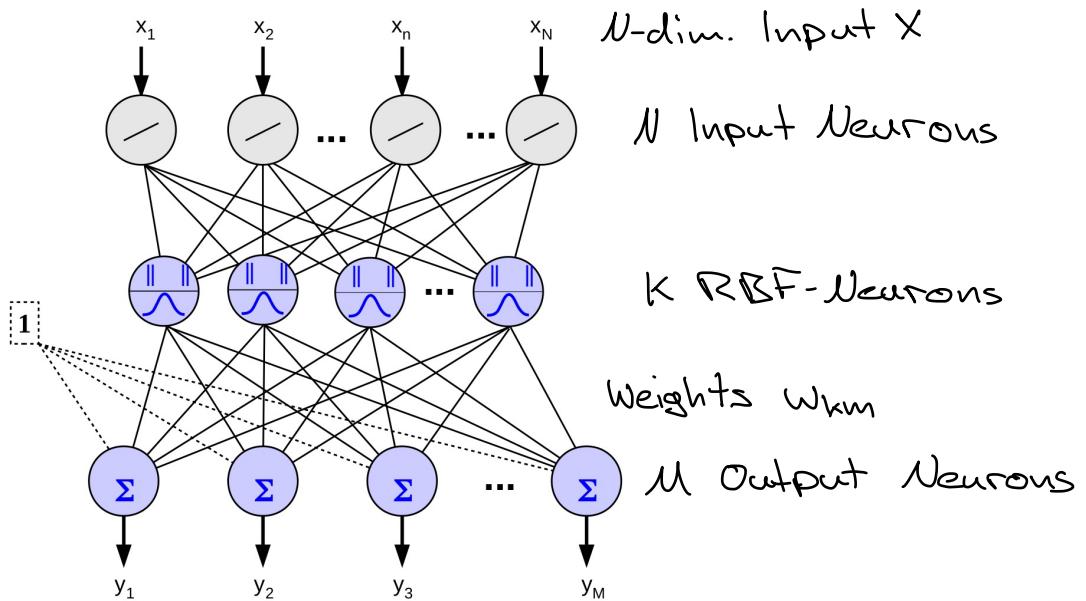
## Open Loop Training

The feedback connections are cut open and the context neurons are artificially set to the teacher values. For Elman networks this technique is more complex, since there are no teacher values for the hidden neurons. Training is performed using Backpropagation



# Radial Basis Function Networks

## Structure:



## RBF Neuron:

Defined by center vector  $c_k$

1. Calculate distance between input and center:  $d_k = \|x - c_k\|$

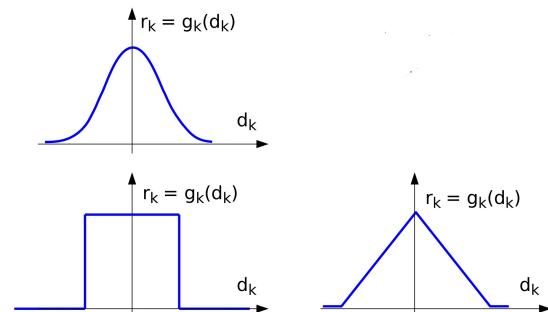
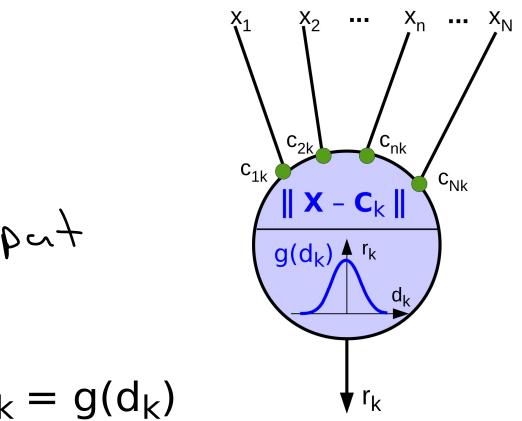
$$d_k = \|x - c_k\|$$

2. Calculate radial function  $g_k$ :  $r_k = g(d_k)$   
→ nonlinear, uni-modal function

Radial Basis Functions:

Gaussian Bell: 
$$r_k = e^{-\frac{d_k^2}{2s_k^2}} = e^{-\frac{\|x - c_k\|^2}{2s_k^2}}$$

with parameter  $s_k$



## RBF Layer:

Accumulation of K RBF neurons

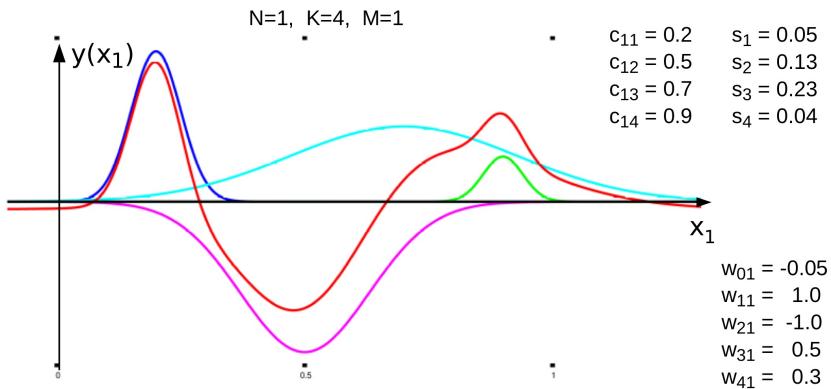
## Output Layer

M neurons that calculate the weighted sum of the RBF neurons output:

$$y_m = \sum_{k=1}^K w_{km} \cdot r_k$$

## Functionality:

### Universal Function Approximator:



## Adjustable Parameters

$c_{nk}$ : centers  $\rightarrow$  Mean of Gaussian

$s_k$ : deviation  $\rightarrow$  width

$w_{km}$ : weight  $\rightarrow$  height

## Setting the RBF neurons

Completely uninformed: Cover input space uniformly

### Input Data driven:

Centers are chosen w.r.t. the distribution of input data:

- Subset of training data
- Statistical Analysis
- Clustering
- Learning

### Error driven:

Centers are adjusted w.r.t. an error function using the teacher values and gradient descent.

## Weight Adjustment

### Simplifications

1. The RBF layer is fixed, fixed  $c_k$  and  $s_k$
2. Type of solution is split in three cases:
  - $P < K$ : not interesting
  - $P = K$ : Interpolation
  - $P > K$ : Approximation, Regression
3. Output size  $M = 1$

### $P = K$ : Interpolation

Teacher values  $\hat{Y}_m$  can be matched exactly by the output

### Linear Equation System

$$\begin{aligned} {}^1\hat{y}_m &\stackrel{!}{=} {}^1y_m(w_{km}) = \sum_{k=1}^K w_{km} \cdot {}^1r_k \\ {}^2\hat{y}_m &\stackrel{!}{=} {}^2y_m(w_{km}) = \sum_{k=1}^K w_{km} \cdot {}^2r_k \quad \Leftrightarrow \\ &\vdots \\ {}^P\hat{y}_m &\stackrel{!}{=} {}^P y_m(w_{km}) = \sum_{k=1}^K w_{km} \cdot {}^P r_k \end{aligned}$$

$$\boxed{\begin{aligned} \hat{Y} &\stackrel{!}{=} \underline{Y} = \underline{R} \cdot \underline{W} \\ \underline{R}^{-1} \cdot \hat{Y} &= \underline{R}^{-1} \underline{R} \cdot \underline{W} = \underline{1} \cdot \underline{W} \\ \underline{R}^{-1} \cdot \underline{Y} &= \underline{W} \end{aligned}}$$

$\Rightarrow$  Weights can be derived directly analytically

### $P > K$ : Approximation/Regression

$R$  is not quadratic and thus not invertible

### Approximation:

Use Moore-Penrose-Pseudo-Inverse:  $\underline{\underline{R}}^+ = (\underline{\underline{R}}^T \cdot \underline{\underline{R}})^{-1} \cdot \underline{\underline{R}}^T$

$$\Rightarrow \boxed{\underline{\underline{R}}^+ \cdot \underline{\hat{Y}} = \underline{W}}$$

### Regression

Learn the weights:

$$\Delta w_{km} = \eta \cdot \delta_m \cdot {}^P r_k = \eta \cdot ({}^P \hat{y}_m - {}^P y_m) \cdot {}^P r_k$$

## Learning Vector Quantization

**Structure:** Neurons structured in a grid like the SOM

**Learning Rule:**

### 0: Initialize SOM

Define: Task, Architecture, Patterns, Learning parameters

Initialize center vectors:

1. completely random
2. subset of training patterns
3. initialize by prior knowledge

1: Choose stimulus and apply to SOM

2: Calculate response:  $r_k = \|\mathbf{C}_k - {}^p\mathbf{X}\|_2$

3: Determine winner

Winner takes it all: Neuron with the closest distance

$$\|\mathbf{C}_i - {}^p\mathbf{X}\|_2 \leq \|\mathbf{C}_j - {}^p\mathbf{X}\|_2 \quad \text{for } i, j = 1 \dots K$$

4: Calculate weight changes  $\Delta C$

**LVQ 1:**  $\Delta \mathbf{C}_i = \eta(t) * (+1) * ({}^p\mathbf{X} - \mathbf{C}_i)$  if  $\mathbf{Y}_i = {}^p\hat{\mathbf{Y}}$  Pulled towards, if correct

$\Delta \mathbf{C}_i = \eta(t) * (-1) * ({}^p\mathbf{X} - \mathbf{C}_i)$  if  $\mathbf{Y}_i \neq {}^p\hat{\mathbf{Y}}$  Pushed away, if not

$\Delta \mathbf{C}_k = 0$  if  $k \neq i$  only winner neuron is trained

**LVQ 2.** Consider the two best neurons  $\mathbf{Y}_i$  and  $\mathbf{Y}_j$

Learn, if: 1. Classes are different:  $\mathbf{Y}_i \neq \mathbf{Y}_j$

2. One is correct:  $\mathbf{Y}_i = \hat{\mathbf{Y}} \vee \mathbf{Y}_j = \hat{\mathbf{Y}}$

3. Stimulus is close to the border of  $\mathbf{Y}_i$  and  $\mathbf{Y}_j$

$$\Delta \mathbf{C}_i = \eta(t) * (+1) * ({}^p\mathbf{X} - \mathbf{C}_i)$$

$$\Delta \mathbf{C}_j = \eta(t) * (-1) * ({}^p\mathbf{X} - \mathbf{C}_j)$$

else  $\Delta \mathbf{C}_k = 0$

**LVQ 3:** Addition to LVQ 2.1

If both classes are correct:

$\alpha$ : learning factor

$$\Delta \mathbf{C}_i = \eta(t) * \alpha * ({}^p\mathbf{X} - \mathbf{C}_i)$$

$$\Delta \mathbf{C}_j = \eta(t) * \alpha * ({}^p\mathbf{X} - \mathbf{C}_j)$$

5: Update step:  $\text{new} \mathbf{C}_k = \text{old} \mathbf{C}_k + \Delta \mathbf{C}_k$

6: Termination criterium, else continue learning (step 1)

Criteriums: Pre-defined performance, Convergence, Time

7: Labeling

Give each neuron a semantic meaning (label / place coding)

## Support Vector Machines

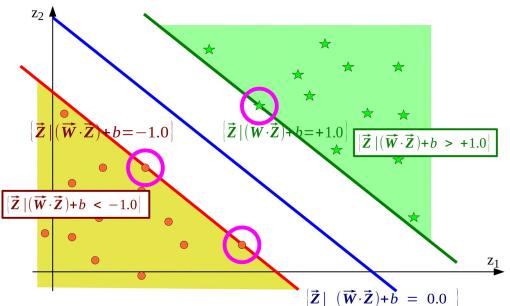
**Task:** Binary classification from linear separable data

**Optimal linear separating hyperplanes**

Hyperplane laying in the middle of the Maximum Margin corridor

$$\text{Width: } \left( \frac{\vec{W}}{\|\vec{W}\|} \cdot (\vec{Z}_* - \vec{Z}_o) \right) = \frac{2}{\|\vec{W}\|}$$

This corridor is spanned by multiple support vectors



**Goal of Learning:** Find the optimal linear separating hyperplane

Let  $X_* \rightarrow 1$  and  $X_o \rightarrow -1$

1. Maximize the margin

$$\text{Maximize the margin } \frac{2}{\|\vec{W}\|} \Leftrightarrow \text{Minimize } \|\vec{W}\| \Leftrightarrow \text{Minimize } \frac{1}{2} \|\vec{W}\|^2$$

2. Classify all patterns correctly

$$\text{Decision function: } f_{W,b}(\vec{Z}) = \text{sign}((\vec{W} \cdot \vec{Z}) + b)$$

$$\text{For all training values } X_i: \begin{cases} f(X_*) = +1 & \Leftrightarrow y_i ((\vec{W} \cdot \vec{X}_i) + b) \geq +1.0 \\ f(X_o) = -1 & \Leftrightarrow y_i ((\vec{W} \cdot \vec{X}_i) + b) - 1 \geq 0.0 \end{cases}$$

$$\text{Lagrange Formalism: } L(\vec{W}, b, \vec{\alpha}) = \frac{1}{2} \|\vec{W}\|^2 - \sum_{i=1}^L \alpha_i (y_i ((\vec{W} \cdot \vec{X}_i) + b) - 1)$$

$\alpha_i$ : Lagrange-Multipliers

Minimize Lagrange function  $L$  w.r.t.  $W$  and  $b$

Maximize Lagrange function  $L$  w.r.t.  $\alpha_i$

$$\text{Optimal solution at } \frac{\partial L(W, b, \alpha)}{\partial W} = \frac{\partial L(W, b, \alpha)}{\partial b} = \frac{\partial L(W, b, \alpha)}{\partial \alpha} = 0$$

**Optimization Problem:**

$$\text{Dual Representation: } D(\vec{\alpha}) = \sum_{i=1}^L \alpha_i - \frac{1}{2} \sum_{i=1}^L \sum_{j=1}^L \alpha_i \alpha_j y_i y_j (\vec{X}_i \cdot \vec{X}_j)$$

Maximize w.r.t. the lagrange multipliers  $\alpha_i$  with  $\sum_{i=1}^L \alpha_i y_i = 0$  &  $\alpha_i \geq 0$

$$\text{Derivation: } \max_{\alpha} \min_{w,b} L(w, b, \alpha) = \frac{1}{2} w \cdot w - \sum_j \alpha_j [(w \cdot x_j + b) y_j - 1]$$

$$\alpha_j \geq 0, \forall j$$

$$\left. \begin{array}{l} \frac{\partial L}{\partial w} = 0 \Rightarrow w = \sum_j \alpha_j y_j x_j \\ \frac{\partial L}{\partial b} = 0 \Rightarrow \sum_j \alpha_j y_j = 0 \end{array} \right\} \Rightarrow \begin{array}{l} \text{maximize}_{\alpha} \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j x_i \cdot x_j \\ \sum_i \alpha_i y_i = 0 \\ \alpha_i \geq 0 \end{array}$$

$$\text{Decision Function: } f_{\vec{\alpha}, b}(\vec{Z}) = \text{sgn} \left( \sum_{i=1}^L \alpha_i y_i (\vec{X}_i \cdot \vec{Z}) + b \right)$$

→ Scalar product between training patterns  $X$  and input  $Z$   
 →  $w$  got eliminated

# Embedding in Higher Feature Space $F$

## Motivation:

If data is not linearly separable in problem space  $G$ , it is linearly separable in a higher dimensional feature space  $F$  with a high probability

**Proof:** Thm. (Cover, 1964): the number of linearly separable dichotomies of  $n$  points in general position in  $\mathbb{R}^d$  is

$$C(n, d) = 2 \sum_{k=0}^d \binom{n-1}{k}$$

Find mapping  $\Phi: G \xrightarrow{\Phi} F$   $\vec{Z} \xrightarrow{\Phi} \Phi(\vec{Z})$   $\vec{X}_i \xrightarrow{\Phi} \Phi(\vec{X}_i)$

$$\Rightarrow D(\vec{\alpha}) = \sum_{i=1}^L \alpha_i - \frac{1}{2} \sum_{i=1}^L \sum_{j=1}^L \alpha_i \alpha_j y_i y_j (\Phi(\vec{X}_i) \cdot \Phi(\vec{X}_j))$$

$$f_{\vec{\alpha}, b}(\vec{Z}) = \text{sgn} \left( \sum_{i=1}^L \alpha_i y_i (\Phi(\vec{X}_i) \cdot \Phi(\vec{Z})) + b \right)$$

→ Embedding in high dimensional spaces is highly inefficient

$\Phi$	problem space	feature space
space	$G = \mathbb{R}^N$	$F = \mathbb{R}^M$
input vectors	$\vec{Z} \in G$	$\Phi(\vec{Z}) \in F$
training vectors	$\vec{X}_i \in G$	$\Phi(\vec{X}_i) \in F$
scalar product	$(\vec{X}_i \cdot \vec{Z})$	$(\Phi(\vec{X}_i) \cdot \Phi(\vec{Z}))$

Kernel Trick: Optimization problem:  $D(\vec{\alpha}) = \sum_{i=1}^L \alpha_i - \frac{1}{2} \sum_{i=1}^L \sum_{j=1}^L \alpha_i \alpha_j y_i y_j k_{\Phi}(\vec{X}_i, \vec{X}_j)$

Decision Function:  $f_{\vec{\alpha}, b}(\vec{Z}) = \text{sgn} \left( \sum_{i=1}^L \alpha_i y_i k_{\Phi}(\vec{X}_i, \vec{Z}) + b \right)$

Instead of embedding both vectors in  $F$ , we can calculate the scalar product of both vectors in  $F$  without the mapping through Mercer kernels:

$$|\Phi(\vec{X}_i) \cdot \Phi(\vec{Z})| = k_{\Phi}(\vec{X}_i, \vec{Z})$$

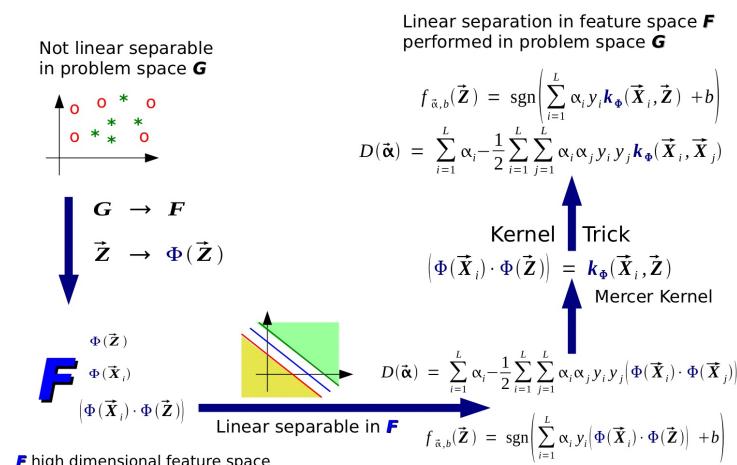
## Mercer Kernels:

Linear Kernel:  $k(\vec{X}_i, \vec{Z}) = (\vec{X}_i \cdot \vec{Z})$

Polynomial Kernel:  $k(\vec{X}_i, \vec{Z}) = (\vec{X}_i \cdot \vec{Z})^d$   
 $k(\vec{X}_i, \vec{Z}) = (s(\vec{X}_i \cdot \vec{Z}) + q)^d$

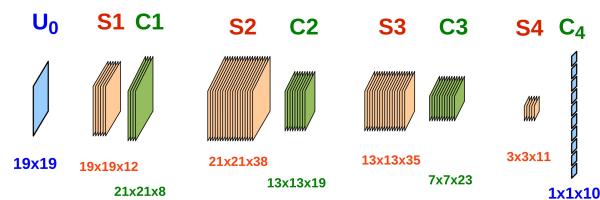
Gaussian Kernel:  $k(\vec{X}_i, \vec{Z}) = \exp(-\|\vec{X}_i - \vec{Z}\|^2/s)$

Tanh Kernel:  $k(\vec{X}_i, \vec{Z}) = \tanh(s(\vec{X}_i \cdot \vec{Z}) + q)$

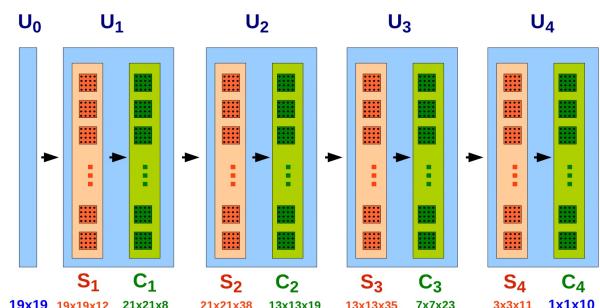


# Neocognitron

Structure:



Modules - S+C Layers - S+C Planes - Neurons



Feed forward sequence of modules ( $U_0, U_1, U_2, \dots$ ):

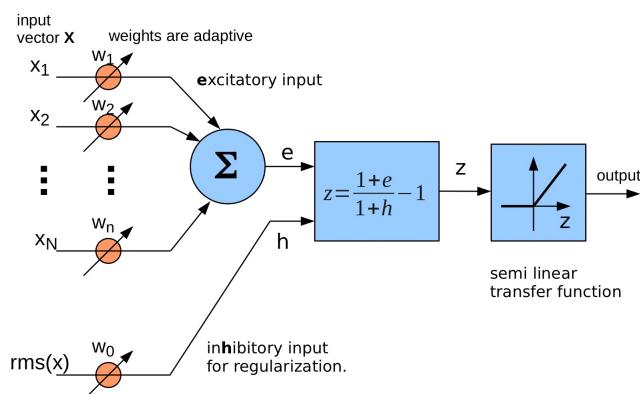
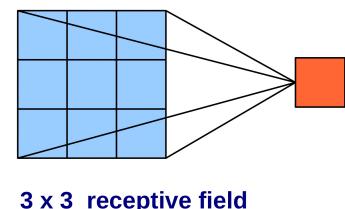
Each module consists of two layers (**S1, C1**) (**S2, C2**, ...)

Each layer is composed by several planes **S-planes** and **C-planes**

Each plane is a 2D arrangement of neurons

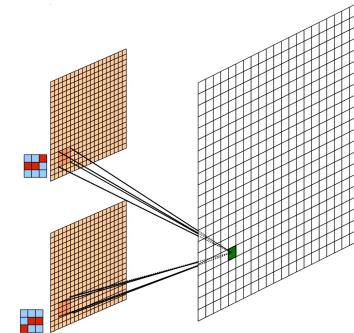
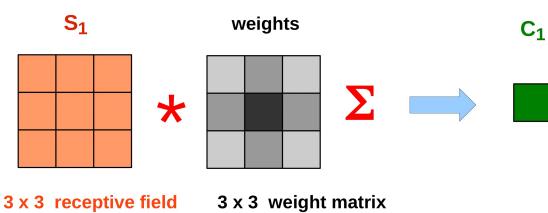
Feature Extractor Cell (S-Cells):

Cell with a  $n \times n$  receptive field that is active if a specific pattern is present.

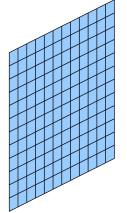


Complex Cells (C-Cells):

Cell consisting of a fixed weight matrix



## Functionality



Module U<sub>0</sub>: Input Layer consisting of a photoreceptor array

Size: 19x19

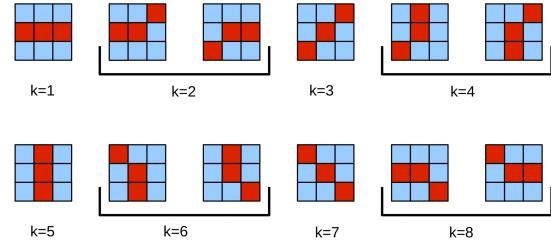
Cells: U<sub>0</sub>(n), with 2D coordinates n = (n<sub>x</sub>, n<sub>y</sub>)

Module U<sub>1</sub>:

S<sub>1</sub>: Extracts simple line features

C<sub>1</sub>: Implements a more robust representation

Features: 12 × 19 × 19 → 8 × 19 × 19

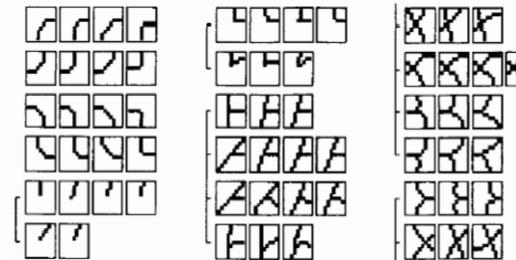


Module U<sub>2</sub>:

S<sub>2</sub>: Extracts more complex features like corners, interactions and curves

C<sub>2</sub>: Robustness (fuzziness)

Features: 38 × 21 × 21 → 19 × 13 × 13

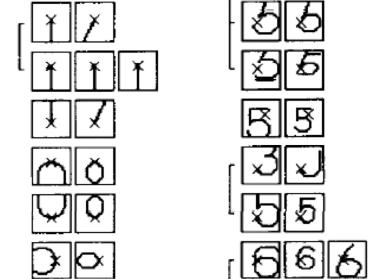


Module U<sub>3</sub>:

S<sub>3</sub>: More complex Features

C<sub>3</sub>: Robustness (fuzziness)

Features: 35 × 13 × 13 → 23 × 7 × 7

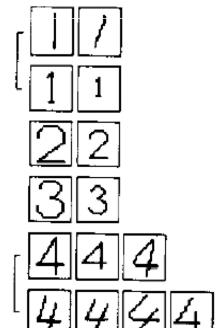


Module U<sub>4</sub>:

S<sub>4</sub>: High-level features

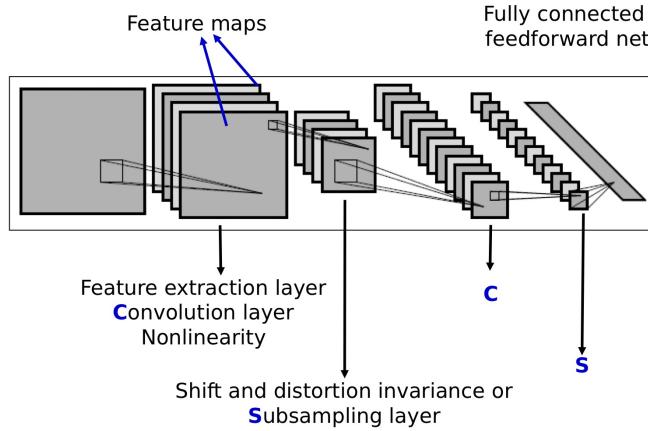
C<sub>4</sub>: Mapping to a one out of 10 coding

Features: 11 × 3 × 3 → 10 × 1 × 1



# Convolutional Neural Networks

## Structure:



## Convolutional Layers:

Extracts features using convolution with a weight matrix and an additional nonlinear mapping.

Features: Represented through the weights

→ weights need to be identical within a single convolutional layer

Channels: Several features can be extracted parallel in a layer through an additional channel in the kernel

Nonlinear Activation: Typically ReLU, Softplus

Weight Sharing: Accumulate weight changes

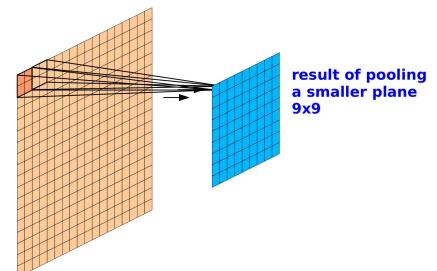
## Subsampling/Pooling Layers

Accumulate positional close features in a single result to implement a more positional independent, sparse representation. Also enables capturing features on a larger scale.

**Max Pooling:** Take max value

0.9	0.3
0.7	0.5

} 0.9  
Max pooling



**Average Pooling:** Take average value

0.9	0.3
0.7	0.5

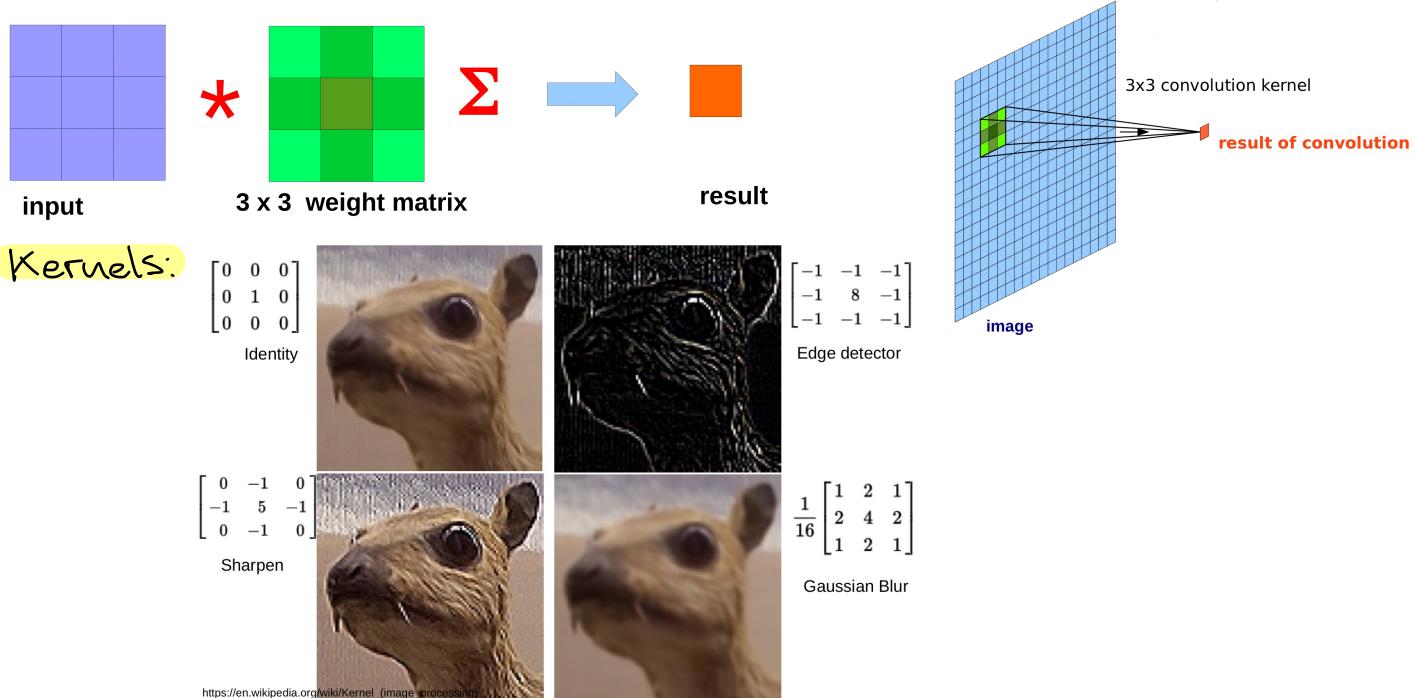
} 0.6  
Average pooling

feature plane  
18 x 18

## Decision: Feed-forward Net

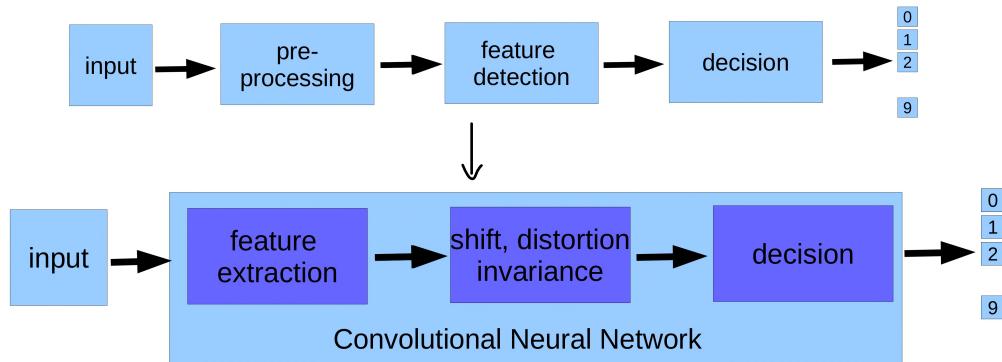
A final MLP is used to use the extracted features for a classification task or accumulate them in a feature descriptor

## Discrete Convolution: Weighted summation



## Task:

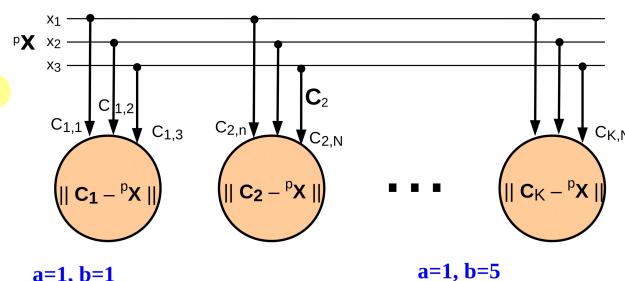
Classification tasks, predominantly in the computer vision.



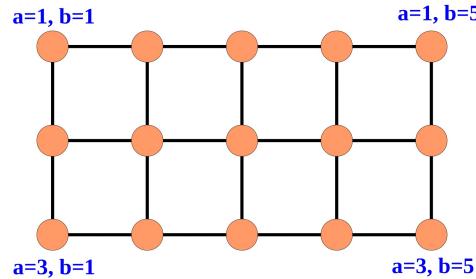
### 3. Unsupervised Learning

#### Self Organizing Feature Maps

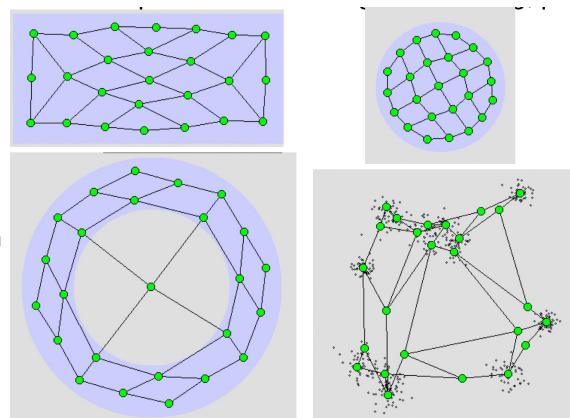
Structure:



Neurons:



Grid:



Task:

- Representation of data by a distribution/grid
  - topology preserving
  - distribution preserving
- Dimensionality reduction
- Clustering

Functionality:

1. Choose stimulus (input vector  $X$ ) and apply to SOM
2. Calculate response (Comparison to center vectors):  $\|c_k - pX\|_2$
3. Determine winner (closest neuron, WTA)
4. Interpret results:
  - labeling: Neurons are labeled
  - place-encoding: Position in the grid reveals interpretation

## Learning Rule:

### O: Initialize SOM

Define: Task, Architecture, Patterns, Learning parameters

Initialize center vectors:

1. completely random
2. subset of training patterns
3. initialize by prior knowledge

1: Choose stimulus and apply to SOM

2: Calculate response:  $r_k = \|\mathbf{C}_k - {}^p\mathbf{X}\|_2$

3: Determine winner

Winner takes it all: Neuron with the closest distance

$$\|\mathbf{C}_i - {}^p\mathbf{X}\|_2 \leq \|\mathbf{C}_j - {}^p\mathbf{X}\|_2 \quad \text{for } i,j = 1 \dots K$$

4: Calculate weight changes  $\Delta C$

$$\Delta \mathbf{C}_j = \eta(t) * h(\text{dist}(i,j), t) * ({}^p\mathbf{X} - \mathbf{C}_j)$$

$\text{dist}(i,j)$ : Distance function defined on the grid

$h(\text{dist}, t)$ : Neighborhood function

5: Update step:  $\text{new } \mathbf{C}_k = \text{old } \mathbf{C}_k + \Delta \mathbf{C}_k$

6: Termination criterium, else continue learning (step 1)

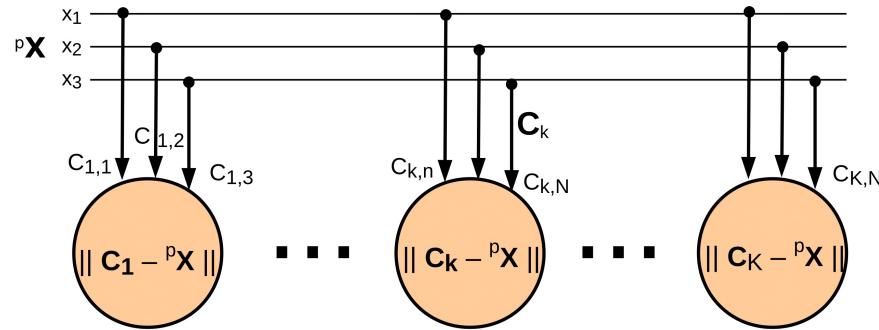
Criteriums: Pre-defined performance, Convergence, Time

7: Labeling

Give each neuron a semantic meaning (label / place coding)

# Neural Gas

## Structure:



Line: Neurons are structured in a sorted list w.r.t the respons  
→ a mesh can be formed after training

## Task:

- Representation of data by a distribution / grid
  - topology preserving
  - distribution Preserving
- Dimensionality Reduction
- Clustering

## Functionality:

1. Choose stimulus (input vector  $X$ ) and apply to SOM
2. Calculate response (Comparison to center vectors):  $\|C_k - {}^p X\|_2$
3. Determine winner (closest neuron, WTA)
4. Interpret results:
  - labeling: Neurons are labeled
  - place-encoding: Position in the grid reveals interpretation

## Comparison to SOM:

- Can preserve topology of more complex structures

## Learning Rule:

### 0: Initialize Neural Gas

Define: Task, Architecture, Patterns, Learning parameters

Initialize center vectors:

1. completely random
2. subset of training patterns
3. initialize by prior knowledge

1: Choose stimulus and apply to Neural Gas

2: Calculate response:  $r_k = \|\mathbf{C}_k - {}^p\mathbf{X}\|_2$

3: Determine winner

Winner takes it all: Neuron with the closest distance

$$\|\mathbf{C}_i - {}^p\mathbf{X}\|_2 \leq \|\mathbf{C}_j - {}^p\mathbf{X}\|_2 \text{ for } i, j = 1 \dots K$$

4: Calculate weight changes  $\Delta C$

$$\Delta \mathbf{C}_j = \eta(t) * h(\text{dist}(i,j), t) * ({}^p\mathbf{X} - \mathbf{C}_j)$$

$\text{dist}(i,j)$ : Neurons are structured in a sorted list w.r.t their response.

Distance function is the position in the list

$h(\text{dist}, +)$ : Neighborhood function

→ Gaussian, linear, hyperbolic

→ Should include all neurons

5: Update step:  $\text{new } \mathbf{C}_k = \text{old } \mathbf{C}_k + \Delta \mathbf{C}_k$

6: Termination criterium, else continue learning (step 1)

Criteriums: Pre-defined performance, Convergence, time

7: Generate a mesh to connect the neurons

- k-nearest neighbor
- Delaunay Triangulation / Voronoi Tesselation
- Statistics from the learning process
- application from prior knowledge

8: Labeling

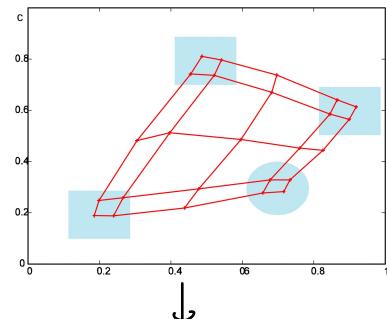
Give each neuron a semantic meaning (label / place coding)

## Multi SOMs/Neural GAS

### Problem SOMs/Neural GAS:

- Clustered data can not be approximated in a topology-preserving manner
- Results in neurons representing data not existing
- Neurons are close that are not close w.r.t. their semantic meaning

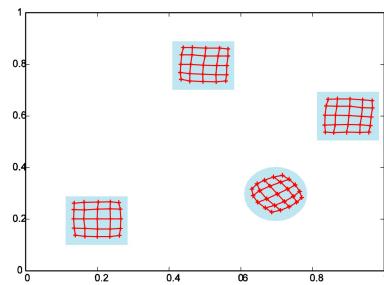
1 SOM:



### Number of SOMs/Neural GAS:

- Determined through prior knowledge
- Learned during training

4 SOM:



### Learning Rule:

Only the SOM/Neural Gas including the winner neuron is trained with their respective learning rule

## K-means Clustering

Step 0: Get some **training data**

Step 1: **Set** the number  $k$  of codebook-vectors

Step 2: **Initialize** the codebook-vectors  $\mathbf{C}_k$

Step 3: **Assign** all patterns to the closest codebook-vector

Step 4: Calculate **centroids** (mean) for each cluster

Step 5: Make the centroids to be the new codebook-vectors  $\mathbf{C}_k$

Step 6: Continue with step 3 **until convergence**

# Regional and Online Learnable Fields (ROLFs)

## Structure:

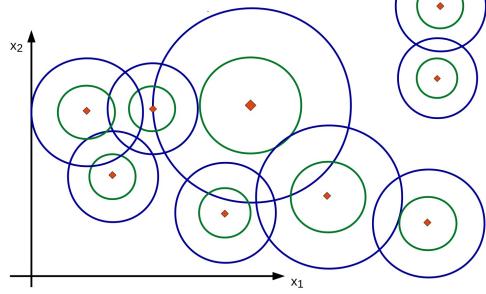
**Neurons:** Center  $C_k$

Individual Size  $\sigma_k$

Perceptive Area with size  $r_k = p \cdot \sigma_k$   
 $\rightarrow p$  is globally defined

Neuron **accepts**, if: Stimulus is in perceptive area

$\rightarrow$  Neurons are not structured in a grid



## Learning Rule:

0: Initialization

Acquire data, set  $p$  and choose  $\sigma$ -strategy

1: FOR EACH Pattern  ${}^p X$

2: IF a neuron accepts

2.1 Calculate distances:  $d = \| {}^p X - C_i \|$

2.1 Adapt neuron  $i$  with closest distance

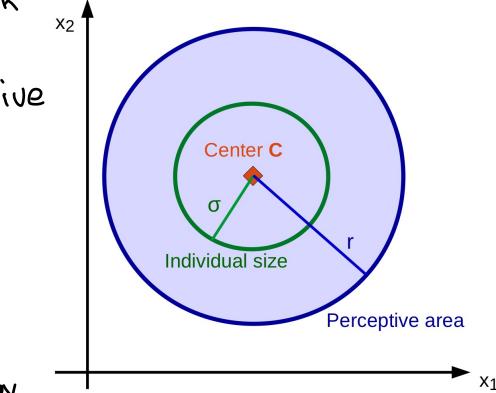
Center  $C_i$ :

$$C_i(t+1) = C_i(t) + \eta_c * ({}^p X - C_i)$$

Size  $\sigma_i$ :

$$\sigma_i(t+1) = \sigma_i(t) + \eta_\sigma * (d - \sigma_i)$$

Perceptive area:  $r_i = p * \sigma_i$



3: ELSE Create a new ROLF neuron

Initialize center  $C_k$ :  $C_{k+1} = {}^p X$

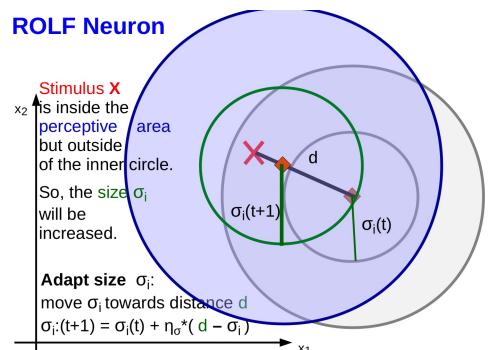
Initialize size  $\sigma_k$ :

Init- $\sigma$ : Use an initial value  $\sigma_{init}$

Minimum- $\sigma$ : Use the minimum of all  $\sigma_k$

Maximum- $\sigma$ : Use the maximum of all  $\sigma_k$

Mean- $\sigma$ : Use the mean of all  $\sigma_k$

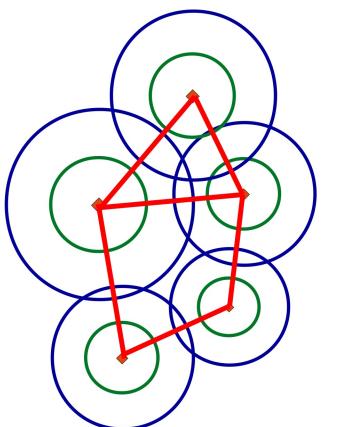


4. In rare cases: repeat for all patterns

5. Build a mesh

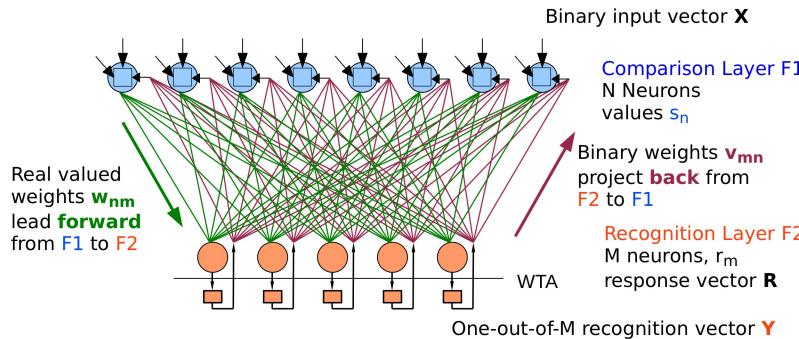
Two neurons are connected, if the perceptive areas overlap:

$$\| C_j - C_k \| \leq (\sigma_j + \sigma_k) * p$$



# Adaptive Resonance Theory (ART-1)

## Structure:



F1: Comparison Layer } Short-term memory  
F2: Recognition Layer }

Forward weight matrix  $W$ : Connecting F1 to F2  
→ real valued

Backward weight matrix  $V$ : Connecting F2 to F1  
→ binary valued

} Long-term memory

## Tasks:

- Pattern recognition
- Content-addressable memory
- Unsupervised Learning
- Clustering

## Functionality:

### F1 comparison layer:

Check, if the prototypic pattern from the recognition layer fits the input vector  $X$

#### Inputs:

1. Binary input vector:  $X = \{x_1, x_2, x_3, \dots, x_n, \dots, x_N\}$

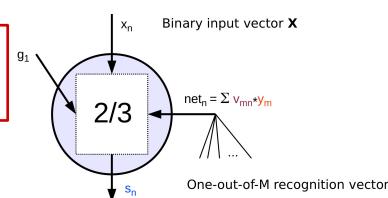
2. Pattern from the recognition layer:

$$net_n = \sum_m v_{mn} * y_m = v_{jn} * y_j = v_{jn}$$

3. Control sign gain  $g_1$ , identical for all neurons

#### Neuron output:

$$s_n = \begin{cases} 1 & \text{if } (x_n * net_n == 1) \text{ OR } (x_n * g_1 == 1) \text{ OR } (net_n * g_1 == 1) \\ 0 & \text{else} \end{cases}$$



### F2 recognition layer:

Recognize a class and output a one-out-of-M output

#### Neuron output:

1. Response:  $r_m = \sum_n w_{nm} * s_n$

2. Winner-takes-it-all:  $y_j = \begin{cases} 1 & \text{if } r_j \geq r_m \\ 0 & \text{else} \end{cases}$

ART-1 Functionality: Structured in 4 phases controlled by g1 & g2

1. Recognition: Mapping input  $X$  to output  $Y$

$$X \rightarrow F1 \rightarrow S \rightarrow W \rightarrow F2 \rightarrow \text{WTA} \rightarrow Y$$

2. Comparison: Mapping output  $Y$  back to  $F1$

$$Y \rightarrow V \rightarrow F1 \xrightarrow{\substack{\text{two-third rule} \\ \curvearrowright}} S'$$

3. Search: Find adequate class in recognition layer

Pattern Recognition:

1. Present input  $X$
2. Network converges to a stable state
3. Active neuron in  $F2$  corresponds to class

Content Addressable Memory:

1. Set one  $F2$  neuron to 1 (corresponding to the class)
2. Corresponding pattern is created in  $F1$

4. Training: Set weights of forward and backward weight matrix

Forward weight matrix  $V$ :

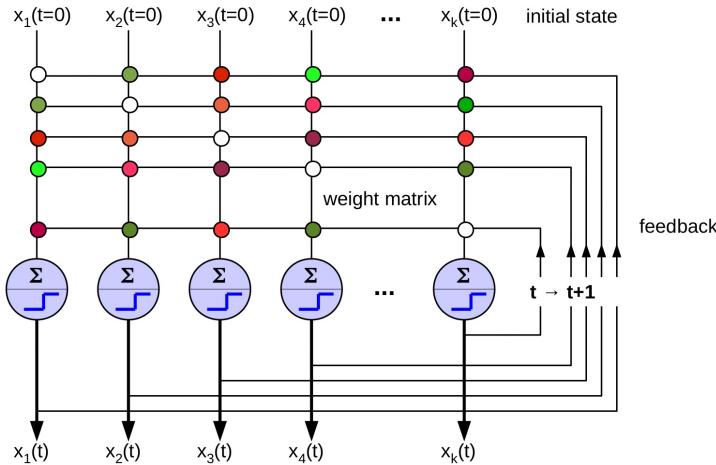
1. Hebbian learning rule: WTA results  $Y$  as teacher or iterative adaption using response
2. Setting directly to WTA based values

Backward weight matrix  $V$ :

$X$  can be used as a teacher, weights can be set to the desired  $S$

# Hopfield Networks

## Structure:



**Neurons:** binary states

**Input:** Weighted sum of the input and feedback connections from every neuron

**Output:** Step function

**Weights:**  $k \times k$  weight matrix

## Tasks:

- content addressable memory
- associative memory
- solve optimization problem

Network state at time t	Weight matrix	
$X(t)$	$\begin{matrix} 0 & +1 & +1 & -1 & +1 & +1 \\ +1 & 0 & -1 & +1 & -1 & -1 \\ +1 & -1 & 0 & +1 & -1 & -1 \\ -1 & +1 & +1 & 0 & -3 & +1 \\ +1 & -1 & -1 & -3 & 0 & +3 \\ -3 & -1 & -1 & +1 & -1 & 0 \\ +1 & -1 & -1 & -3 & +3 & -1 \end{matrix}$	
$\begin{matrix} 1 & +1 & +1 & -1 & +1 & +1 \\ -1 & 0 & -1 & +1 & -1 & -1 \\ -1 & -1 & 0 & +1 & -1 & -1 \\ -1 & +1 & +1 & 0 & -3 & +1 \\ -1 & -1 & -1 & -3 & 0 & +3 \\ -3 & -1 & -1 & +1 & -1 & 0 \\ +1 & -1 & -1 & -3 & +3 & -1 \end{matrix}$	$=$	$\begin{matrix} -4 & 0 & -2 & +8 & -2 & +4 & -2 \end{matrix}$

The threshold function  $f(z)$ , generates the new binary state

## Functionality:

State of the neurons  $X(t)$  is the processed information

1. Initialize with initial pattern  $X(t=0)$

2. Update  $X(t+1)$  until convergence or specific  $t$  is reached:

**Update rule:**  $x_k(t+1) = f_k \left( \sum_{j=1}^K w_{jk} \cdot x_j(t) \right)$ , with  $f_k(z) = \begin{cases} +1 & \text{if } z > \theta_k \\ -1 & \text{if } z \leq \theta_k \end{cases}$

or  $f_k(z) = \begin{cases} +1 & \text{if } z > \theta_k \\ X_k(t) & \text{if } z = \theta_k \\ -1 & \text{if } z < \theta_k \end{cases}$

## Synchronous Update:

All states  $x(t+1)$  are updated in one step

## Asynchronous Update:

In a timestep only one neuron is updated. The following updates depend on the already updated neurons.

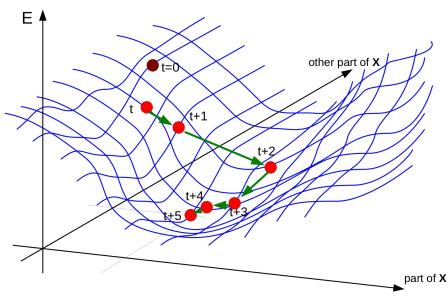
The order of updates is random

→ used in Hopfield networks

## Dynamics:

Energy function:  $E = -\frac{1}{2} \sum_i^K \sum_j^K w_{ij} x_i x_j + \sum_k^K \theta_k x_k$

→ decrease, thus network converges to a stable state after a while



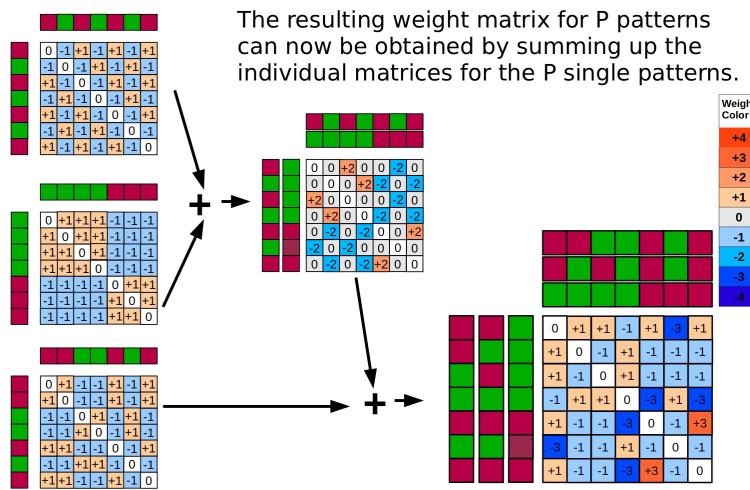
## Learning:

Learning is done by storing the patterns in the weight matrix

Weight  $x_{ik}$ :

$$w_{ik} = \sum p_{X_i} * p_{X_k}$$

$$\begin{cases} 1 & , \text{if } x_i = x_k \\ -1 & , \text{else} \end{cases}$$



Maximal number of patterns:  $Q = \alpha K$ , with  $\alpha \approx 0.138 \dots \approx 0.14$

## Modes:

### Autoassociator Mode

Associate a pattern with itself. Weights are calculated w.r.t. the components of the same pattern

$$w_{jk} = x_j \cdot x_k$$

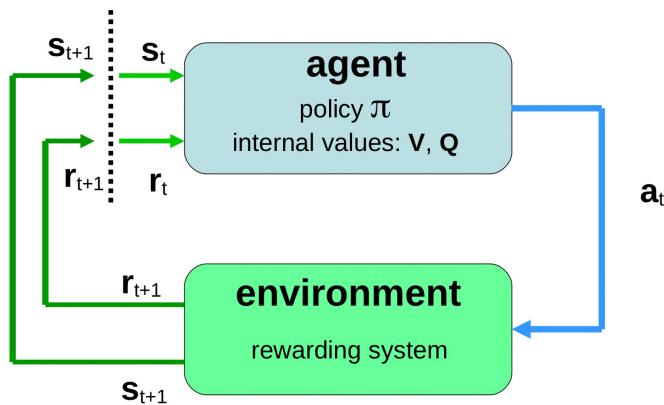
### Heteroassociator

Associate pattern  $p$  with a second pattern  $q$ . Weights are calculated w.r.t. to two patterns.

$$w_{jk} = {}^p x_j \cdot {}^q x_k$$

### 3. Reinforcement Learning

Learning Model:



Agent: Sees state  $s_t$  and decides on action  $a_t$  based on:

Policy  $\pi$ : A mapping from states to actions to maximize the amount of expected reward

State Value Function  $V^\pi(s)$ : 
$$V^\pi(s) = E_\pi[R_t | s_t = s] = E_\pi \left[ \sum_{k=0}^{\infty} \gamma^k r_{t+1+k} | s_t = s \right]$$

Expected reward w.r.t  $\pi$  starting from this state

Action Value Function  $Q^\pi(s, a)$ : 
$$Q^\pi(s, a) = E_\pi[R_t | s_t = s, a_t = a] = E_\pi \left[ \sum_{k=0}^{\infty} \gamma^k r_{t+1+k} | s_t = s, a_t = a \right]$$

Expected reward when in state  $s$  taking action  $a$

Environment: Receives action  $a_t$  and responds with a new state  $s_{t+1}$ . Additionally a feedback as a reward following the action is provided

State Space: Space containing the possible discrete states

Action Space: Contains all possible actions performable by the agent

Rewarding System: Produces scalar rewards

Goal: Find a policy  $\pi$  that maximizes the reward and (hopefully) solves a given task optimal

## Rewarding

True Return:  $R_t^* = r_{t+1}^* + r_{t+2}^* + r_{t+3}^* + r_{t+4}^* + \dots + r_T^*$

- accumulated future rewards
- in general, not obtainable

Expected Return:  $R_t = r_{t+1} + r_{t+2} + r_{t+3} + r_{t+4} + \dots + r_T$

Temporal Horizon:  
 $T = \infty$ : infinite temporal horizon  
 $T < \infty$ : future can be split into finite temporal episodes

Discounting:  $R_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \gamma^3 r_{t+4} + \dots$

$$R_t = \sum_{k=0}^T \gamma^k r_{(t+1+k)} \quad 0 \leq \gamma \leq 1.0$$

$\gamma$ : Discounting Factor

Weigh impact of future rewards on the current return

Special Rewards:

### 1. Pure delayed reward

Only the last reward  $r_T$  of the episode, i.e. the result of the episode, holds a value, rest is zero, e.g.:

$$r_t = r_T = 0.0 \text{ if } t < T \text{ or the outcome is draw, stalemate}$$

$$r_T = +1.0 \text{ if game has been won}$$

$$r_T = -1.0 \text{ if game has been lost}$$

### 2. Pure negative reward

The agent is punished with a negative reward for every step he takes, e.g.:

$$r_t = -1.0 \text{ if } t < T, \text{ pure negative reward during the episode}$$

$$r_T = 0.0 \text{ no reward at the end of the episode}$$

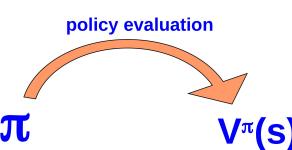
### 3. Avoidance Strategy

Almost all actions are not rewarded, except for a few with a negative reward that shall be avoided, e.g.:

$$r_t = -1.0 \text{ if } s_t \in \{ !, \ominus, +, \times, \div \} \text{ for selected situations.}$$

$$r_t = 0.0 \text{ else, in the rest of the cases (majority).}$$

## Policy Evaluation



Determining the value functions given a specific policy  $\pi$

**Behavior Policy:** Policy used for evaluation to explore the environment and determining the value functions

**Prediction:** Policy evaluation w.r.t the state value function  $V$

## Monte Carlo Policy Evaluation

Estimate the state value function by taking random actions. The value changes as soon as a return is available, i.e. a terminal state is reached.

### First-Visit MC evaluation:

Averages all returns after the first time state  $s$  was visited

### Every-Visit MC evaluation:

Averages all returns for all visits of state  $s$

Constant- $\alpha$ : 
$$V(s_t) \leftarrow V(s_t) + \alpha [R_t - V(s_t)]$$

## Algorithm

0: Initialize policy  $\pi$  and state value function

1: UNTIL termination criterion:

2: Generate episode using  $\pi$

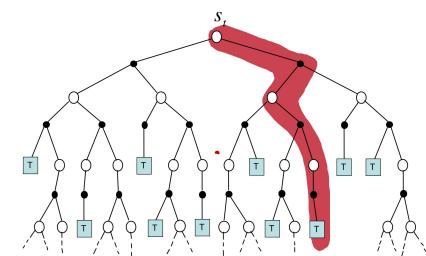
3: FOR EACH state  $s$  in episode

4:  $R \leftarrow$  Return following  $s$

5: Append  $R$  to Returns( $s$ )

6:  $V(s) \leftarrow \text{Average}(\text{Returns}(s))$

Drawback: Returns are only available at the end of every episode



## Temporal Difference Policy Evaluation

Instead of using the true return, approximate it through the obtained reward  $r_{t+1}$  and an estimate of the forthcoming rewards  $\gamma V(s_{t+1})$ .

Temporal Difference: 
$$V(s_t) \leftarrow V(s_t) + \alpha [r_{t+1} + \gamma V(s_{t+1}) - V(s_t)]$$

Derivation: 
$$R_t = \sum_{k=0}^{\infty} \gamma^k r_{t+1+k} = r_{t+1} + \gamma \sum_{k=0}^{\infty} \gamma^k r_{t+1+k+1} = r_{t+1} + \gamma V(s_{t+1})$$

Control: Policy evaluation w.r.t. the action value function  $Q$

### On-policy Methods:

Evaluate/or improve the policy that is used to make decision

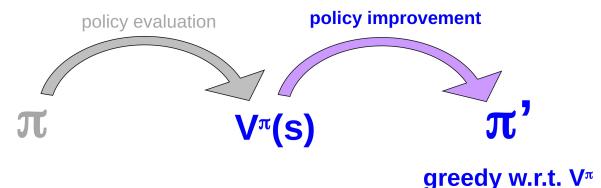
SARSA: 
$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$

### Off-policy Methods:

Separate policies to generate behavior (Behavior Policy) and to evaluate or improve (Estimation Policy)

Q-Learning: 
$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [r_{t+1} + \gamma \max_a Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$

### Policy Improvement



Improve the policy given the state value function in a greedy manner w.r.t  $V^\pi$ .

Improved Policy:  $\pi' \rightarrow \text{greedy}(V) \Leftrightarrow \pi(s) = \arg \max_a Q(s, a)$

$$\begin{aligned} Q^{\pi_k}(s, \pi_{k+1}(s)) &= Q^\pi(s, \arg \max_a Q^{\pi_k}(s, a)) \\ &= \max_a Q^{\pi_k}(s, a) \\ &\geq Q^{\pi_k}(s, \pi_k(s)) \\ &= V^{\pi_k}(s) \end{aligned}$$

⇒ The new policy is at least as good as the old one

### General Policy Iteration

Iterate until convergence:

1. Policy Evaluation
2. Policy Improvement

Convergence: optimal policy and state value function

